

# Estimating an HPC Facility's Capacity For Accommodating Realtime Workflows

Eashan Adhikarla<sup>1</sup>, Brian Austin<sup>2</sup>  
<sup>1</sup>National Energy Research Scientific Computing (NERSC), <sup>2</sup>Lehigh University



## Introduction

Supercomputers are valuable, highly contended resources that are typically managed by batch queue systems.

A growing number of workloads (including tightly-integrated experimental facilities and emergency decision making) require response-times that are not compatible with typical queue wait-times.

Any policy to accommodate real-time requirements will impact system utilization and disrupt other users.

This project estimates the consequences of three potential real-time scheduling policies by analyzing job logs from NERSC's Cori Supercomputer.

## Methodology


We obtained from NERSC, a log of the node count, start time, end time and memory footprint for all jobs that ran on Cori's KNL nodes during 2018. We randomly sampled 50-timepoints that correspond to the submission times for a set of hypothetical realtime jobs. The job logs were then analyzed to compute median wait-time, utilization-loss or memory availability for each of the proposed scheduling policies described below.

### Cori and its Workload at a Glance

|               |                               |
|---------------|-------------------------------|
| Compute Nodes | 9688 (KNL), 2388 (Haswell)    |
| Jobs per day  | 60,000                        |
| Nodes per Job | 1-9688, Avg.= 716 hours       |
| Time per Job  | <1min – 8 days, Avg.=19 hours |


## Results

**Draining:** Upon submission of a real-time job, the compute nodes are drained by allowing currently running batch jobs to complete normally and prohibiting batch jobs from starting until the real-time job begins.





**Job Cancellation:** Upon submission of a real-time job, a fraction of the currently running batch jobs are cancelled to allow the real-time job to start immediately. Jobs are selected for cancellation to minimize waste of computational resources.

This policy is extremely disruptive to users whose jobs are cancelled.



**Job Pausing:** Upon submission of a real-time job, a fraction of the currently running batch jobs are paused (using, for example, the 'nice' command) to allow the real-time job to start immediately. The paused jobs maintain their memory footprint, so jobs are selected for pausing to maximize the memory available to the real-time job.



## Conclusions

Because HPC resources are finite, no scheduling policy can provide sufficient compute elasticity to run arbitrarily large jobs on demand. All the policies evaluated here "passes the buck" in a different direction. System-draining delays the start of the real-time job and squanders HPC resources with idle nodes. Job-cancellation gives sufficient priority to real-time jobs, but at the expense of severe disruption to other batch users whose jobs are cancelled. "Pausing" batch jobs is a promising option that allows real-time jobs to start immediately and avoids wasting system resources but requires more elaborate changes to the scheduler and may increase memory pressure, especially for high-concurrency real-time jobs.

## Future Work

- Introduce supervised machine learning model to the set of policies so that the model can make internally decisions on which policy to be allotted to the incoming job.
- Improve these results by accounting Backfill.

## References

- Time-sharing redux for large-scale HPC systems : [ieeexplore.ieee.org/document/7828392](http://ieeexplore.ieee.org/document/7828392)
- NERSC-2014 Workload Analysis : [http://portal.nersc.gov/project/mpccc/baustin/nersc2014workloadanalysis\\_v1.1](http://portal.nersc.gov/project/mpccc/baustin/nersc2014workloadanalysis_v1.1)

